

HOMework 2

CMU 10-707: DEEP LEARNING (FALL 2017)

<https://piazza.com/cmu/fall2017/10707>

OUT: Oct 9, 2017

DUE: Oct 23, 2017 11:59 pm

TAs: Otilia Stretcu, Shunyuan Zhang

Instructions

Hand in answers to all questions below. For Problem 5, the goal of your write-up is to document the experiments you have done and your main findings, so be sure to explain the results. Be concise and to the point – do not write long paragraphs or only vaguely explain results.

- The answers to all questions should be in pdf form (please use \LaTeX). Do not hand-write your submission- we reserve the right to take off points or not accept hand-written submissions.
- Please include a README file with instructions on how to execute your code. Your code should contain the implementation for the RBM, autoencoder and denoising autoencoder. Please create your own implementations and do not use any toolboxes.
- Package your code and README document using `zip` or `tar.gz` in a file called `10707-A2-*yourandrewid*. [zip|tar.gz]`.
- Submit your PDF write-up to the Gradescope assignment “Homework 2” and your packaged code to the Gradescope assignment “Code for Homework 2.”

Problem 1 (10 pts)

Consider a simple convolutional neural network (CNN) applied to a $d \times d$ image. This CNN contains a single convolutional layer with a kernel of size $k \times k$, stride s pixels, and 2 feature maps, followed by a single $p \times p$ subsampling (max pooling) layer, followed by a softmax layer of size m . Derive the backpropagation algorithm with respect to the adjustable parameters of this network and discuss how the standard backpropagation algorithm must be modified when evaluating the derivatives of an error function with respect to the adjustable parameters in the network.

Problem 2 (10 pts)

Consider a directed graphical model with K random variables. By marginalizing out the variables in order, show that the representation

$$p(\mathbf{x}) = \prod_{k=1}^K p(x_k | \mathbf{pa}_k) \quad (1)$$

for the joint distribution of a directed graph is correctly normalized, provided each of the conditional distributions is normalized.



Figure 1: An example of undirected graphical model called a Restricted Boltzmann Machine.

Problem 3 (10pts)

In class, we considered a bipartite undirected graphical model, called Restricted Boltzmann Machine (RBM), shown in Figure 1. An RBM contains a set of observed and hidden random variables, where binary visible variables $v_i \in \{0, 1\}$, $i = 1, \dots, D$, are connected to binary hidden variables $h_j \in \{0, 1\}$, $j = 1, \dots, P$. The energy of the joint configuration is given by:

$$E(\mathbf{v}, \mathbf{h}; \theta) = - \sum_{i=1}^D \sum_{j=1}^P W_{ij} v_i h_j - \sum_{i=1}^D v_i b_i - \sum_{j=1}^P h_j a_j. \quad (2)$$

where $\theta = \{W, \mathbf{a}, \mathbf{b}\}$ denote model parameters, with \mathbf{a} and \mathbf{b} representing the bias terms, and $\mathbf{v} = \{v_1, \dots, v_D\}$, $\mathbf{h} = \{h_1, \dots, h_P\}$. The probability of the joint configuration is given by the Boltzmann distribution:

$$P_{\theta}(\mathbf{v}, \mathbf{h}) = \frac{1}{\mathcal{Z}} \exp(-E(\mathbf{v}, \mathbf{h}; \theta)), \quad (3)$$

where $\mathcal{Z} = \sum_{\mathbf{v}, \mathbf{h}} \exp(-E(\mathbf{v}, \mathbf{h}; \theta))$ denotes the normalizing constant. Derive that conditioned on observed variables \mathbf{v} , the distribution over the hidden variables factorizes (so the hidden variables become independent conditioned on \mathbf{v}):

$$P_{\theta}(h_1, \dots, h_P | \mathbf{v}) = \prod_{j=1}^P p_{\theta}(h_j | \mathbf{v}), \quad (4)$$

where

$$P_{\theta}(h_j = 1 | \mathbf{v}) = \frac{1}{1 + \exp(-\sum_{i=1}^D W_{ij} v_i - a_j)}. \quad (5)$$

Note: inferring the states of the hidden variables conditioned on the observed data is easy. This represents a major advantage of this model, which is successfully used in many domains, ranging from collaborative filtering to speech recognition.

Problem 4 (10 pts)

Consider the use of iterated conditional modes (ICM) to minimize the energy function (that we considered in class for image denoising example) given by:

$$E(\mathbf{x}, \mathbf{y}) = h \sum_i x_i - \beta \sum_{i,j} x_i x_j - \eta \sum_i x_i y_i, \quad (6)$$

where $x_i \in \{-1, 1\}$ is a binary variable denoting the state of pixel i in the unknown noise-free image, i and j are indices of neighboring pixels, $y_i \in \{-1, 1\}$ denotes the corresponding value of pixel i in the observed noisy image, and h, β and η are positive constants. The joint distribution is defined as:

$$p(\mathbf{x}, \mathbf{y}) = \frac{1}{Z} \exp(-E(\mathbf{x}, \mathbf{y})) \quad (7)$$

- (5 pts) Write down an expression for the difference in the values of the energy associated with the two states of a particular variable x_j , with all other variables held fixed, and show that it depends only on quantities that are local to x_j in the graph.
- (5 pts) Consider a particular case of the energy function above in which the coefficients $\beta = h = 0$. Show that the most probable configuration of the latent variables is given by $x_i = y_i$ for all i .

Problem 5 (60 pts)

For this question you will write your own implementation of the Contrastive Divergence algorithm for training RBMs, and compare it to an Autoencoder model. Please do not use any toolboxes. We recommend that you use MATLAB or Python, but you are welcome to use any other programming language if you wish.

The goal is to build a generative model of images of 10 handwritten digits of “zero”, “one”, ..., “nine”. The images are 28 by 28 in size (MNIST dataset), which we will be represented as a vector \mathbf{x} of dimension 784 by listing all the pixel values in raster scan order. The labels t are 0,1,2,...,9 corresponding to 10 classes as written in the image. There are 3000 training cases, containing 300 examples of each of 10 classes, 1000 validation (100 examples of each of 10 classes), and 3000 test cases (300 examples of each of 10 classes). they can be found in the file digitstrain.txt, digitsvalid.txt and digitstest.txt:

<http://www.cs.cmu.edu/~rsalakhu/10707/assignments.html>

Format of the data: digitstrain.txt contains 3000 lines. Each line contains 785 numbers (comma delimited): the first 784 real-valued numbers correspond to the 784 pixel values, and the last number denotes the class label: 0 corresponds to digit 0, 1 corresponds to digit 1, etc. digitsvalid.txt and digitstest.txt contain 1000 and 3000 lines and use the same format as above.

Contrastive Divergence (CD), Autoencoders

Implement the CD algorithm for training an RBM model. Implement both an autoencoder and a denoising autoencoder model.

a) Basic generalization [20 points]

Train an RBM model with 100 hidden units, starting with CD with $k = 1$ step. For initialization use samples from a normal distribution with mean 0 and standard deviation 0.1. Choose a reasonable learning rate (e.g. 0.1 or 0.01). Use your own judgement to decide on the stopping criteria (i.e number of epochs or convergence criteria).

As a proxy for monitoring the progress, plot the average training cross-entropy reconstruction error on the y-axis vs. the epoch number (x-axis). On the same figure, plot the average validation cross-entropy error function.

Examine the plots of training error and validation error. How does the network's performance differ on the training set versus the validation set during learning? Visualize the learned W as 100 28x28 images (plot all filters as one image, as we have seen in class). Do the learned features exhibit any structure?

b) Number of CD steps [5 points]

Try learning an RBM model with CD with $k = 5$, $k = 10$, and $k = 20$ steps. Describe the effect of this modification on the convergence properties, and the generalization of the network.

c) Sampling from the RBM model [5 points]

To qualitatively test the model performance, initialize 100 Gibbs chains with random configurations of the visible variables, and run the Gibbs sampler for 1000 steps. Display the 100 sampled images. Do they look like handwritten digits?

d) Unsupervised Learning as pretraining [5 points]

Use the weights you learned in question **a)** to initialize a 1-layer neural network with 100 hidden units. Train the resulting neural network to classify the 10 digits, as done in the first assignment. Does this pretraining help compared to training the same neural network initialized with random weights in terms of classification accuracy? Describe your findings.

e) Autoencoder [5 points]

Instead of training an RBM model, train an autoencoder model with 100 sigmoid hidden units. For initializing the autoencoder, use samples from a normal distribution with mean 0 and standard deviation 0.1. Does the model learn useful filters? Visualize the learned weights. Similar to RBMs, use the learned autoencoder weights to initialize a 1-layer neural network with 100 hidden units. Train the resulting neural network to classify the 10 digits. How does it compare to pretraining with RBMs?

f) Denoising Autoencoder [10 points]

Train a denoising autoencoder by using drop-out on the inputs. We recommend a drop-out rate of 10%, but feel free to explore other rates. Visualize the learned weights. Use the learned weights to initialize a 1-layer neural network with 100 hidden units. Train the resulting neural network to classify the 10 digits. How does it compare to pretraining with standard autoencoders and RBMs in terms of classification accuracy?

g) Number of hidden units [10 points]

Try training RBMs, autoencoders and denoising autoencoders with 50, 100, 200, and 500 hidden units. Describe the effect of this modification on the convergence properties, and the generalization of the network.